

D2: Dual-Source $q(\mathbf{F})$ Update Under $\eta = \mathbf{YF}\tilde{\beta}$

Cluster B (Cox-on-YF Reformulation) — Derivation D2 (Primary)

Cross-reference: `derivations/Modular_UpdateDerivations/` $q(L)$ section ($n \leftrightarrow p$ symmetry)

Andrew Walther

2026-04-29

Contents

1	Context and Goal	2
2	Model Recap	2
3	Factor-wise ELBO for $q(f_{jk})$	2
3.1	Genomics Contribution	3
3.2	Survival Contribution	3
3.3	Combined EBNM Inputs	4
4	Scale Imbalance Audit	4
4.1	Original Problem (for reference)	4
4.2	Cluster B Scaling for $q(\mathbf{F})$	4
5	Identifiability Under $\eta = \mathbf{Z}_F\tilde{\beta}$	5
6	Train/Test Consistency	5
7	Code Implications	5
8	Verification Checklist	6

1 Context and Goal

Under the Cluster B reformulation $\eta_i = (\mathbf{y}_i \mathbf{F}) \tilde{\beta}$, the factor matrix $\mathbf{F} \in \mathbb{R}^{p \times K}$ appears in *both*:

1. the genomics likelihood $\mathbf{Y} \approx \mathbf{L} \mathbf{F}^\top$, and
2. the Cox PH likelihood through $\eta_i = \sum_k (\sum_j y_{ij} f_{jk}) \tilde{\beta}_k$.

This is analogous to how $\mathbf{L}$ had a dual role in the original model. The derivation follows the same structure as the original $q(\mathbf{L})$ derivation (`Modular_UpdateDerivations/ q(L)` section), with the n - and p -dimensional roles swapped.

Goal of D2. Derive the factor-wise ELBO for $q(f_{jk})$, identify the dual-source EBNM inputs $A_{F,j}$ and $B_{F,j}$ for both sources, and establish that the scale imbalance between survival and genomics is resolved under this formulation.

2 Model Recap

Generative model.

$$y_{ij} \mid \mathbf{l}_i, \mathbf{f}_j, \tau_j \sim \mathcal{N}(\mathbf{l}_i^\top \mathbf{f}_j, \tau_j^{-1}), \quad h_i(t) = h_0(t) e^{\eta_i}, \quad \eta_i = \sum_k \underbrace{\left(\sum_j y_{ij} f_{jk} \right)}_{(z_F)_{ik}} \tilde{\beta}_k.$$

Notation.

- $\mathbf{Z}_F = \mathbf{Y} \mathbf{E}[\mathbf{F}] \in \mathbb{R}^{n \times K}$: projection scores held fixed per outer CAVI iteration.
- $w_i = W_{ii}$: negative diagonal Hessian weights from the Cox Taylor expansion.
- $z_{\text{no-}k,i} = z_i - \sum_{k' \neq k} (z_F)_{ik'} \mathbb{E}[\tilde{\beta}_{k'}]$: partial working response with factor k removed.
- $R_k = \mathbf{Y} - \sum_{k' \neq k} \mathbb{E}[\mathbf{l}_{k'}] \mathbb{E}[\mathbf{f}_{k'}]^\top$: partial genomics residual.

CAVI update frequency. $\mathbf{Z}_F$ is recomputed once per outer CAVI iteration (before the Cox Taylor expansion). Within the k -loop, $(z_F)_{ik} = \sum_j y_{ij} \mathbb{E}[f_{jk}]$ — so after $\mathbb{E}[\mathbf{f}_k]$ is updated, $\mathbf{Z}_F$ is stale. However, the standard Gauss-Seidel contract holds $\mathbf{Z}_F$ fixed for the duration of the inner k -loop; it is refreshed at the start of the next outer iteration. This matches how η is treated in the original model.

3 Factor-wise ELBO for $q(f_{jk})$

In Gauss-Seidel CAVI, we update $q(\mathbf{f}_k)$ holding all other quantities fixed.

3.1 Genomics Contribution

Expanding the squared error for gene j , factor k :

$$\begin{aligned}\mathbb{E}_q[\log p(\mathbf{Y} \mid \mathbf{L}, \mathbf{F}, \boldsymbol{\tau})] \Big|_{j,k} &= -\frac{\tau_j}{2} \sum_i \mathbb{E}_q[(r_{k,ij} - l_{ik} f_{jk})^2] + \text{const} \\ &= -\frac{\tau_j}{2} \left[\mathbb{E}_q[f_{jk}^2] \sum_i \mathbb{E}_q[l_{ik}^2] - 2 \mathbb{E}_q[f_{jk}] \sum_i \mathbb{E}_q[r_{k,ij}] \mathbb{E}_q[l_{ik}] \right] + \text{const.}\end{aligned}\quad (1)$$

Reading off the coefficient of $\mathbb{E}_q[f_{jk}^2]$ and $\mathbb{E}_q[f_{jk}]$:

$$A_{F,\text{gen},j} = \tau_j \sum_i \mathbb{E}_q[l_{ik}^2], \quad (2)$$

$$B_{F,\text{gen},j} = \tau_j \sum_i \mathbb{E}_q[r_{k,ij}] \mathbb{E}_q[l_{ik}] = \tau_j [\mathbf{R}_k^\top \mathbb{E}[\mathbf{L}_k]]_j. \quad (3)$$

3.2 Survival Contribution

The Cox PH likelihood enters via the Taylor-linearised surrogate. The relevant expected log-likelihood (up to constants in $q(\mathbf{f}_k)$) is

$$\mathbb{E}_q[\log p(\mathbf{t}, \boldsymbol{\delta} \mid \mathbf{Y}, \mathbf{F}, \tilde{\boldsymbol{\beta}})] \approx \sum_i \left[u_i \eta_i - \frac{1}{2} w_i \eta_i^2 \right] + \text{const}, \quad (4)$$

where u_i is the score residual and the term in η_i^2 is the Taylor quadratic.

Substituting $\eta_i = \sum_{k'} (z_F)_{ik'} \mathbb{E}[\tilde{\beta}_{k'}]$ and isolating the contribution from f_{jk} (through $(z_F)_{ik} = \sum_{j'} y_{ij'} f_{j'k}$):

Isolating f_{jk} in the survival term

Write $\eta_i = y_{ij} f_{jk} \mathbb{E}[\tilde{\beta}_k] + \tilde{\eta}_{\text{no-}j,i}$ where $\tilde{\eta}_{\text{no-}j,i}$ collects all terms that do not depend on f_{jk} . Then (4) contributes, for gene j , factor k :

$$\begin{aligned}& -\frac{1}{2} \sum_i w_i \mathbb{E}_q[\eta_i^2] \Big|_{f_{jk}} \\ &= -\frac{1}{2} \mathbb{E}_q[f_{jk}^2] \mathbb{E}[\tilde{\beta}_k^2] \sum_i w_i y_{ij}^2 + \mathbb{E}_q[f_{jk}] \mathbb{E}[\tilde{\beta}_k] \sum_i w_i z_{\text{no-}k,i} y_{ij} + \text{const},\end{aligned}\quad (5)$$

where $z_{\text{no-}k,i}$ is the partial working response ($\eta_{\text{no-}k,i} + u_i/w_i - \sum_{j'} y_{ij'} \mathbb{E}[f_{j'k}] \mathbb{E}[\tilde{\beta}_k]$ simplifies to $z_i - (z_F)_{\cdot,-k} \mathbb{E}[\beta_{-k}]$, cancelling the contribution of factor k).

Reading off the coefficients:

$$A_{F,\text{surv},j} = \mathbb{E}[\tilde{\beta}_k^2] \sum_i w_i y_{ij}^2 = \mathbb{E}[\tilde{\beta}_k^2] [\text{diag}(\mathbf{Y}^\top \text{diag}(\mathbf{w}) \mathbf{Y})]_j, \quad (6)$$

$$B_{F,\text{surv},j} = \mathbb{E}[\tilde{\beta}_k] \sum_i w_i z_{\text{no-}k,i} y_{ij} = \mathbb{E}[\tilde{\beta}_k] [\mathbf{Y}^\top (\mathbf{w} \circ \mathbf{z}_{\text{no-}k})]_j. \quad (7)$$

3.3 Combined EBNM Inputs

Mixing the two sources with weight $\alpha \in [0, 1]$:

$$A_{F,j} = (1 - \alpha) A_{F,\text{gen},j} + \alpha A_{F,\text{surv},j}, \quad (8)$$

$$B_{F,j} = (1 - \alpha) B_{F,\text{gen},j} + \alpha B_{F,\text{surv},j}. \quad (9)$$

EBNM call for $q(f_k)$ under Cluster B

Inputs (p-vectors):

$$x_{F,j} = \frac{B_{F,j}}{A_{F,j}}, \quad s_{F,j} = \frac{1}{\sqrt{A_{F,j}}}.$$

Note on τ_j cancellation: In the pure-genomics limit ($\alpha = 0$), $x_{F,j} = B_{F,\text{gen},j}/A_{F,\text{gen},j}$, and τ_j cancels in $x_{F,j}$ (divides numerator and denominator equally) but *not* in $s_{F,j}$ (which has $A_{F,\text{gen},j}^{-1/2} \propto \tau_j^{-1/2}$). This is the same τ -cancellation property as in the original $q(\mathbf{F})$ derivation. When $\alpha > 0$, τ_j does not cancel globally (survival and genomics A terms have different τ dependencies), but numerically the expression in (8) is well-defined.

Prior family: point-exponential (or as configured in `prior_LF`)

4 Scale Imbalance Audit

A key question: does the Cluster B formulation fix the scale imbalance that plagued the original $q(\mathbf{L})$ update?

4.1 Original Problem (for reference)

In the original model ($\eta = \mathbf{L}\beta$), the $q(\mathbf{L})$ precision was:

$$A_{L,\text{gen}} = \sum_j \tau_j \mathbb{E}[f_{jk}^2] \approx p \bar{\tau} \mathbb{E}[f^2] \quad \text{vs.} \quad A_{L,\text{surv},i} = \lambda w_i \mathbb{E}[\beta_k^2] \approx O(1).$$

The ratio $A_{L,\text{surv}}/A_{L,\text{gen}} \approx 1/(p \bar{\tau}) \approx 1/2000$ for PDAC-scale data — survival cannot compete.

4.2 Cluster B Scaling for $q(\mathbf{F})$

Genomics precision.

$$A_{F,\text{gen},j} = \tau_j \sum_i \mathbb{E}[l_{ik}^2].$$

This sums over n patients ($n \approx 200\text{--}500$). Typical values: $\tau_j \approx 0.1\text{--}10$, $\sum_i \mathbb{E}[l_{ik}^2] \approx n \sigma_l^2 \approx n \cdot 0.01$. So $A_{F,\text{gen},j} \approx n \tau_j \sigma_l^2$.

Survival precision.

$$A_{F,\text{surv},j} = \mathbb{E}[\tilde{\beta}_k^2] \sum_i w_i y_{ij}^2.$$

Typical values: $w_i \approx 0.05\text{--}0.2$ (Cox weight), $y_{ij}^2 \approx 1$ (standardised expression), $\mathbb{E}[\tilde{\beta}_k^2] \approx 0.01\text{--}0.1$. So $A_{F,\text{surv},j} \approx n \bar{w} \bar{y}^2 \mathbb{E}[\tilde{\beta}^2] \approx n \cdot 0.1 \cdot 1 \cdot 0.05$.

Ratio.

$$\frac{A_{F,\text{surv},j}}{A_{F,\text{gen},j}} \approx \frac{\bar{w} \bar{y}^2 \mathbb{E}[\tilde{\beta}^2]}{\tau_j \sigma_l^2} \approx \frac{0.1 \cdot 1 \cdot 0.05}{1 \cdot 0.01} = 0.5.$$

Scale imbalance resolved

The n factors cancel in the ratio — both $A_{F,\text{gen}}$ and $A_{F,\text{surv}}$ scale with n . The ratio is $O(1)$ rather than $O(1/p)$. Survival and genomics contributions to $q(\mathbf{F})$ are now comparable in magnitude, which is the structural fix Cluster B was designed to achieve.

5 Identifiability Under $\eta = \mathbf{Z}_F \tilde{\beta}$

Scaling ambiguity

Under $\eta_i = \sum_k (\sum_j y_{ij} f_{jk}) \tilde{\beta}_k$, there is a scaling ambiguity: scaling $f_{jk} \rightarrow c f_{jk}$ and $\tilde{\beta}_k \rightarrow \tilde{\beta}_k / c$ leaves η_i unchanged. This is equivalent to the $\mathbf{L}$ - β scaling ambiguity in the original model.

The EBNM prior on $q(\mathbf{f}_k)$ breaks this ambiguity by imposing sparsity. Specifically:

- The point-exponential prior shrinks small $|f_{jk}|$ to zero and sets a characteristic scale via the estimated σ_k in the EBNM family. This anchors the scale of f_{jk} .
- Given the anchored scale of $\mathbf{F}$, $\tilde{\beta}_k$ is identified through the EBNM prior on $q(\tilde{\beta}_k)$ (point-normal with EBNM-estimated scale).

Practical consequence. No additional orthogonality or normalization constraint is needed. The Gauss-Seidel updates cycle between $q(\mathbf{F})$ and $q(\tilde{\beta})$, and the joint EBNM shrinkage provides implicit identification. This is the same mechanism as the original model (where $q(\mathbf{L})$ and $q(\beta)$ are jointly identified through EBNM).

6 Train/Test Consistency

Train

Training:

$$\hat{\eta}_i = \sum_k (z_F)_{ik} \mathbb{E}[\tilde{\beta}_k] = (\mathbf{y}_i \mathbb{E}[\mathbf{F}]) \mathbb{E}[\tilde{\beta}].$$

Test:

$$\hat{\eta}_i^{\text{test}} = (\mathbf{y}_i^{\text{test}} \mathbb{E}[\mathbf{F}]) \mathbb{E}[\tilde{\beta}].$$

Identical formula. No matrix inverse, no OLS projection, no special-case logic. The learned $\mathbb{E}[\mathbf{F}]$ and $\mathbb{E}[\tilde{\beta}]$ transfer directly.

7 Code Implications

`code/update_F.R` — `update_F_k()`

- **New arguments:** `w` (n-vector), `EBeta_k` (scalar), `EBeta2_k` (scalar), `YtWY_diag` (p-vector, pre-computed), `YtWz_no_k` (p-vector, pre-computed per k)

- New computation:

```

A_F_surv <- EBeta2_k * YtWY_diag      # p-vector: eq. (5)
B_F_surv <- EBeta_k * YtWz_no_k      # p-vector: eq. (6)
A_F <- (1-alpha)*A_F_gen + alpha*pmax(A_F_surv, A_floor)
B_F <- (1-alpha)*B_F_gen + alpha*B_F_surv

```

- Pre-computation in caller (fit_modular.R):

```

YtWY_diag <- colSums(sweep(Y^2, 1, w, "*"))      # once per outer iter
YtWz_no_k_k <- as.vector(t(Y) %*% (w * z_no_k))  # once per k

```

code/update_F.R — update_F_all() Add w, EBeta, EBeta2, Y, z to signature. Pre-compute YtWY_diag once per call; compute YtWz_no_k per k.

8 Verification Checklist

D2 Verification

1. $A_{F,\text{surv},j}$ formula matches (6): $\mathbb{E}[\tilde{\beta}_k^2] \cdot \sum_i w_i y_{ij}^2$. ✓
2. $B_{F,\text{surv},j}$ formula matches (7): $\mathbb{E}[\tilde{\beta}_k] \cdot \sum_i w_i z_{\text{no-}k,i} y_{ij}$. ✓
3. Scale ratio $A_{F,\text{surv}}/A_{F,\text{gen}} \approx O(1)$ (not $O(1/p)$). ✓
4. τ_j cancels in $x_{F,j}$ in the pure-genomics limit ($\alpha = 0$). ✓
5. When $\alpha = 0$: $A_F = A_{F,\text{gen}}$, $B_F = B_{F,\text{gen}}$ (reduces to original). (to test)
6. When $w=0$: $A_{F,\text{surv}} = 0$, $B_{F,\text{surv}} = 0$. (to test)
7. When $\text{EBeta_k}=0$: $B_{F,\text{surv}} = 0$. (to test)
8. All F tests pass including 6 new dual-source tests. (to implement)