

Contents

update_F.R — Companion Document	2
Table of Contents	2
Overview	2
Mathematical Background	2
Model Structure	2
EBNM Sufficient Statistics for $f_{\cdot}$	3
The Tau Cancellation Property	3
Tau Cancels in x_j (observation)	3
Tau Does NOT Cancel in s_j (noise level)	3
Consequence	4
API Reference	4
<code>update_F_k(Tau, EL_k, EL2_k, R_k, prior_family, A_floor)</code>	4
<code>update_F_all(Y, EL, EL2, EF, EF2, Tau, prior_family, A_floor)</code>	5
Test Suite	6
T1: Mathematical Identity Checks (6 tests)	6
T2: WLS / Non-Informative Limit (2 tests)	6
T3: Known Signal Recovery $K=1$ (3 tests)	6
T4: Multi-Factor $K=5$ (3 tests)	7
T5: Null Factor Shrinkage (2 tests)	7
T6: Tau Differential Shrinkage (3 tests)	7
T7: Numerical Stability (4 tests)	7
T8: R_k and Gauss-Seidel (2 tests)	7
T9: V2.R Consistency (1 test)	7
Demos	8
Demo 1: Tau Cancellation Property	8
Demo 2: Signal Recovery $K=1$	8
Demo 3: Tau-Dependent Shrinkage	8
Demo 4: Sparsity Recovery	9
Demo 5: Multi-Factor $K=5$	9
Computational Notes	9
Scalar Broadcast	9
Shared Dependency: <code>compute_R_k</code>	9
Memory Layout	9
Related Files	10

update_F.R — Companion Document

Module: code/update_F.R **Project:** Supervised Bayesian Matrix Factorization (multiomicsGEP) **Role:** Variational update for the genomics loading matrix F (factor columns $f_{\{.k\}}$)

Table of Contents

1. Overview
 2. Mathematical Background
 3. The Tau Cancellation Property
 4. API Reference
 5. Test Suite
 6. Demos
 7. Related Files
-

Overview

update_F.R implements the CAVI variational update for the factor loading matrix F (shape $p \times K$) in the joint genomics + survival model. F appears **only** in the genomics likelihood and has no role in the Cox proportional hazards survival term — this structural fact drives the module's key mathematical simplification.

The update is a **vector EBNM** (Empirical Bayes Normal Means) problem solved independently for each factor column $f_{\{.k\}}$. Within each column, the p feature-wise posterior means and variances are computed in closed form, then passed to an EBNM solver that applies the chosen prior (default: point-normal) and returns posterior moments.

Depends on: compute_R_k() from code/update_L.R — source that file before this one.

Mathematical Background

Model Structure

The genomics observations follow:

$$Y \approx L F' \quad (n \times p \text{ matrix, } L \text{ is } n \times K, F \text{ is } p \times K)$$

with feature-wise precision $\text{Tau}[j]$ (heteroscedastic noise). The residual for factor k at feature j is:

$$R_k[i,j] = Y[i,j] - \sum_{k' \neq k} E_L[i,k'] * E_F[j,k']$$

This is computed by compute_R_k() and excludes the current factor k (Gauss-Seidel schedule).

EBNM Sufficient Statistics for $f_{\cdot}$

For each feature j , the ELBO contribution from factor k leads to the following normal sufficient statistics:

Quantity	Formula	Shape	Notes
$A_F[j]$	$\text{Tau}[j] * \text{sum}(\text{EL2_k})$	p-vector	$\text{sum}(\text{EL2_k})$ is a scalar broadcast across all j
$B_F[j]$	$\text{Tau}[j] * (\text{t}(\text{R_k}) \% \% \text{EL_k})[j]$	p-vector	Inner product of residual column with L posterior
$x_F[j]$	$B_F[j] / A_F[j]$	p-vector	EBNM observation (“sufficient statistic”)
$s_F[j]$	$1 / \text{sqrt}(A_F[j])$	p-vector	EBNM noise standard deviation

$\text{sum}(\text{EL2_k})$ is the sum of the second moments of the k -th column of L across all n samples:

$$\text{sum_EL2_k} = \text{sum_i } \text{EL2}[i, k] \quad (\text{scalar, includes } \text{Var}_q(l_{ik}) + \text{EL}[i, k]^2)$$

The Tau Cancellation Property

This is the **central mathematical insight** of the module.

Tau Cancels in x_j (observation)

$$\begin{aligned} x_j &= B_F[j] / A_F[j] \\ &= (\text{Tau}[j] * (\text{t}(\text{R_k}) \% \% \text{EL_k})[j]) / (\text{Tau}[j] * \text{sum_EL2_k}) \\ &= (\text{t}(\text{R_k}) \% \% \text{EL_k})[j] / \text{sum_EL2_k} \end{aligned}$$

The $\text{Tau}[j]$ factors in numerator and denominator cancel exactly. The EBNM observation x_j is **identical regardless of feature precision**.

Tau Does NOT Cancel in s_j (noise level)

$$s_j = 1 / \text{sqrt}(A_F[j]) = 1 / \text{sqrt}(\text{Tau}[j] * \text{sum_EL2_k})$$

Here $\text{Tau}[j]$ remains under the square root. A feature with higher precision (larger $\text{Tau}[j]$) gets a **smaller** noise level s_j , which means the prior applies less shrinkage and the posterior mean stays closer to x_j .

Consequence

Features with the same residual signal x_j are differentially shrunk toward zero depending on their noise precision. Clean features (high τ) are shrunk less; noisy features (low τ) are shrunk more. This is the heteroscedastic shrinkage mechanism for F .

τ	x_j	s_j	Shrinkage
1	same	large	strong
1000	same	small ($\sim 1/\sqrt{1000}$) smaller	weak

Demo 1 illustrates this with $\max |x_j \text{ diff}| \approx 0$ and $s_j \text{ ratio} \approx \sqrt{1000}$.

API Reference

update_F_k(τ , EL_k , $EL2_k$, R_k , $prior_family$, A_floor)

Computes the variational posterior for a single factor column $f_{\{k\}}$.

Arguments:

Parameter	Type	Description
τ	numeric vector, length p	Feature-wise noise precisions
EL_k	numeric vector, length n	Posterior means of L column k : $E_q[l_{\{k\}}]$
$EL2_k$	numeric vector, length n	Posterior second moments of L column k : $E_q[l_{\{k\}}^2]$
R_k	numeric matrix, $n \times p$	Residual matrix excluding factor k (from <code>compute_R_k</code>)
$prior_family$	character	EBNM prior family; default "point_normal"
A_floor	numeric	Minimum value for A_F to avoid division by zero; default $1e-10$

Returns: Named list

Field	Type	Description
mean	numeric vector, length p	Posterior mean $E_q[f_{\{k\}}]$
second	numeric vector, length p	Posterior second moment $E_q[f_{\{k\}}^2]$
sd	numeric vector, length p	Posterior standard deviation
A	numeric vector, length p	EBNM precision A_F (after floor)
B	numeric vector, length p	EBNM numerator B_F

Field	Type	Description
x	numeric vector, length p	EBNM sufficient statistic $x_F = B/A$
s	numeric vector, length p	EBNM noise $s_F = 1/\sqrt{A}$
sum_EL2_k	numeric scalar	sum(EL2_k), shared across all features
ebnm_result	list	Raw output from <code>ebnm()</code> call

Key implementation details:

- $A_F = \text{pmax}(\text{Tau} * \text{sum}(\text{EL2_k}), A_{\text{floor}})$ — the floor prevents NaN when Tau is zero
- $B_F = \text{Tau} * (\text{t}(\text{R_k}) \%*\% \text{EL_k})$ — matrix-vector product gives p-vector
- Second moment satisfies: $\text{second}[j] = \text{sd}[j]^2 + \text{mean}[j]^2$

`update_F_all(Y, EL, EL2, EF, EF2, Tau, prior_family, A_floor)`

Gauss-Seidel loop over all K factor columns.

Arguments:

Parameter	Type	Description
Y	numeric matrix, $n \times p$	Observed genomics data
EL	numeric matrix, $n \times K$	Posterior means of L (fully updated from L step)
EL2	numeric matrix, $n \times K$	Posterior second moments of L
EF	numeric matrix, $p \times K$	Current posterior means of F (updated in place)
EF2	numeric matrix, $p \times K$	Current posterior second moments of F
Tau	numeric vector, length p	Feature-wise precisions
prior_family	character	EBNM prior family; default "point_normal"
A_floor	numeric	Floor for A; default $1e-10$

Returns: Named list

Field	Type	Description
EF	numeric matrix, $p \times K$	Updated posterior means of F
EF2	numeric matrix, $p \times K$	Updated posterior second moments of F
details	list of length K	Per-factor output from <code>update_F_k</code>

Important Gauss-Seidel note:

update_F_all passes the **fully updated** EL (not a stale copy) because the L step completes before the F step in the CAVI schedule. Within the F loop, EF_curr is updated column-by-column so that each compute_R_k call sees the most recent estimates of all other F columns. This is the Gauss-Seidel (sequential) schedule, as opposed to a Jacobi (parallel) schedule.

```
for k in 1:K:
  R_k <- compute_R_k(Y, EL, EF_curr, k) # excludes column k
  res <- update_F_k(Tau, EL[,k], EL2[,k], R_k, ...)
  EF_curr[,k] <- res$mean
  EF2_curr[,k] <- res$second
```

Test Suite

File: tests/test_update_F.R Total: **26 tests across 9 groups** — all passing.

T1: Mathematical Identity Checks (6 tests)

Test	What is verified
T1.1	$A_F[j] = \text{Tau}[j] * \text{sum}(\text{EL2_k})$ exactly
T1.2	$B_F[j] = \text{Tau}[j] * (\text{t}(\text{R_k}) \%*\% \text{EL_k})[j]$ exactly
T1.3	x is a p-vector (length p)
T1.4	Tau cancels in x_j : Tau=1 vs Tau=1000 produce identical x_j
T1.5	Tau does NOT cancel in s_j : s_j ratio is significantly different
T1.6	Second moment identity: $\text{second}[j] = \text{sd}[j]^2 + \text{mean}[j]^2$

T2: WLS / Non-Informative Limit (2 tests)

Test	What is verified
T2.1	When $\text{EL2_k} = \text{EL_k}^2$ (zero posterior variance in L), x_j equals the OLS estimator
T2.2	Large Tau drives posterior mean close to the observation x_j

T3: Known Signal Recovery K=1 (3 tests)

Test	What is verified
T3.1	Dense factor: correlation between recovered and true f exceeds 0.7
T3.2	Sparse factor (50/500 nonzero): active features have larger posterior mean than inactive
T3.3	Higher Tau $\square$ better recovery (monotone precision effect)

T4: Multi-Factor K=5 (3 tests)

Test	What is verified
T4.1	Output matrices EF, EF2 have shape $p \times K$
T4.2	All second moments $\geq$ squared means ($EF2 \geq EF^2$) component-wise
T4.3	All outputs are finite (no NaN or Inf)

T5: Null Factor Shrinkage (2 tests)

Test	What is verified
T5.1	Pure-noise R_k produces factor near zero (point-normal shrinks to zero)
T5.2	Zero $EL_k \square A$ hits floor, $B=0 \square EF$ near zero

T6: Tau Differential Shrinkage (3 tests)

Test	What is verified
T6.1	High-tau features are shrunk less than low-tau features (same x_j)
T6.2	Larger $EL2 \square$ larger $A \square$ smaller $s \square$ less shrinkage for high-signal factors
T6.3	Uniform Tau $\square$ all s_j equal (homoscedastic limit, one shared noise level)

T7: Numerical Stability (4 tests)

Test	What is verified
T7.1	All-zero Tau $\square A$ hits floor, no NaN or Inf in output
T7.2	Extreme Tau= $1e10 \square$ no overflow in A, B, x, s
T7.3	Edge case $p=1, n=1 \square$ correct scalar outputs
T7.4	Custom A_{floor} value is respected ($A \geq A_{\text{floor}}$)

T8: R_k and Gauss-Seidel (2 tests)

Test	What is verified
T8.1	R_k correctly uses EL_k (not another factor's column)
T8.2	$update_F_all$ produces finite results with valid second moments across all K

T9: V2.R Consistency (1 test)

Test	What is verified
T9.1	update_F_k output matches the corresponding lines in Supervised_Bayesian_MF_V2.R (lines 334–340) exactly

Demos

File: demos/demo_update_F.R Total: **5 scenarios** — narrative output, no PASS/FAIL.

Demo 1: Tau Cancellation Property

Purpose: Demonstrate the core mathematical property that τ_j cancels in x_j but not in s_j .

Setup: Same Y , L , R_k . Run update_F_k twice — once with $\tau = \text{rep}(1, p)$, once with $\tau = \text{rep}(1000, p)$.

Expected output: - $\max_j |x_j(\tau=1) - x_j(\tau=1000)| \approx 0$ (machine precision) - $s_j(\tau=1) / s_j(\tau=1000) \approx \sqrt{1000} \approx 31.6$ for all j

This illustrates that the EBNM “data” is precision-free, while the EBNM “noise level” encodes the precision.

Demo 2: Signal Recovery $K=1$

Purpose: Confirm that the update recovers a sparse true factor from noisy observations.

Setup: $n=150$, $p=200$, $K=1$. True f_{true} is sparse (30/200 nonzero entries, drawn from $N(0,1)$). $Y = L \% \% t(f_{\text{true}}) + \text{noise}$.

Reported metrics: - Pearson correlation between $EF[,1]$ and f_{true} - Mean posterior mean for active vs inactive features (active should be larger)

Demo 3: Tau-Dependent Shrinkage

Purpose: Show that heteroscedastic precision differentially regularizes features.

Setup: $p=200$ features split into two groups: first 100 with $\tau=10$ (clean), last 100 with $\tau=0.5$ (noisy). Same true signal in both groups.

Reported metrics: - Mean absolute error (MAE) for clean vs noisy features - Posterior SDs for clean vs noisy features - Demonstrates that clean features (high τ) stay closer to truth

Demo 4: Sparsity Recovery

Purpose: Evaluate the point-normal prior's ability to identify which features belong to a sparse factor.

Setup: $n=100$, $p=250$, true f is sparse (40/250 nonzero). Threshold posterior means at 0.1 to classify active/inactive.

Reported metrics (confusion matrix):

	Predicted Active	Predicted Inactive
True Active	TP	FN
True Inactive	FP	TN

Also reports precision = $TP/(TP+FP)$ and recall = $TP/(TP+FN)$.

Demo 5: Multi-Factor $K=5$

Purpose: Validate `update_F_all` on a structured 5-factor problem.

Setup: $n=200$, $p=300$, $K=5$. Each factor has a distinct sparse structure; F is initialized at zero.

Reported metrics: - Column-wise correlations: $\text{cor}(EF[:,k], F_{\text{true}}[:,k])$ for each k - Number of active features per factor (posterior mean > 0.1) - Confirms Gauss-Seidel within-loop updates do not break structure

Computational Notes

Scalar Broadcast

`sum(EL2_k)` is computed once (a single scalar) and broadcast across all p features via element-wise multiplication with Tau . This avoids a $p \times n$ matrix multiply for A_F .

Shared Dependency: `compute_R_k`

`compute_R_k` is defined in `code/update_L.R`. Always source `update_L.R` before `update_F.R`. The residual matrix R_k has shape $n \times p$ and represents the data with all other factors' contributions removed. Because L is fully updated before the F step, the R_k computation here uses the most recent L posterior.

Memory Layout

Matrix	Shape	Update order
EL	$n \times K$	Fixed (updated in L step)
EL2	$n \times K$	Fixed (updated in L step)
EF	$p \times K$	Updated column by column (Gauss-Seidel)

Matrix	Shape	Update order
EF2	$p \times K$	Updated column by column
R_k	$n \times p$	Recomputed fresh for each k

For large p , R_k ($n \times p$) dominates memory. It is not cached between factor updates.

Related Files

File	Relationship
code/update_L.R	Defines <code>compute_R_k()</code> required by this module
code/update_tau.R	Provides <code>Tau</code> used as input here
code/update_beta.R	Beta update (survival); <code>F</code> does not appear there
code/Supervised_Bayesian_MF_M2.R	W2n algorithm; lines 334–340 are the reference for T9.1
tests/test_update_F.R	26 tests, all passing
demos/demo_update_F.R	5 narrative demos
derivations/qF/qF_update_derivation.tex	Full derivation including tau cancellation proof
code/SupervisedMF_Context.m4	AI/code quick-reference for the full R implementation